



Princeton Computer Science Contest – Fall 2021

Problem 8: Optimal Search Trees [HackerRank]

By Ruijie Fang

1 Background: Searching in Databases

You are working for a very large database company and they have a database index consisting of n sorted integer keys $\mathcal{S} := s_1, s_2, \dots, s_n$ with $s_i \leq s_{i+1}$ for $1 \leq i < n$, each mapping to some value. The n keys are some predetermined sparse, big integers, which you don't know, meaning they are non-trivial to hash (if you don't know what this means, don't worry about it). Furthermore, you know that the company workload every day generates a fixed access pattern to the database, resulting in access frequencies $F_1, F_2, \dots, F_n \in \mathbb{Z}$. You can of course use a dynamic search tree, like a splay tree, to answer this problem. However your boss wants you to be **optimal**: Let $d(i)$ denote the depth of the i -th node in the array (\mathcal{S}) in the search tree. Let the root node in the search tree have depth 0 (i.e. it can be accessed at no cost). An optimal search tree is a search tree that minimizes the cost function $\sum_{i=1}^n F_i \cdot d(i)$. Intuitively, this means you want the most frequently accessed values to be toward the top, and least frequently accessed toward the bottom.

2 Definitions

Binar Search Tree: Given a set $\mathcal{S} := \{k_1, k_2, \dots, k_n\}$ of n integer keys with $k_1 \leq k_2 \leq \dots \leq k_n$, a binary search tree T_2 is a tree of n binary search tree nodes. A binary search tree node can be represented as (x, c_1, c_2) where $x \in \mathcal{S}$ is a key in \mathcal{S} uniquely represented in T_2 , c_1 is the leftmost child node of x and c_2 is the rightmost child node of x . Let $S(1, x)$ denote the set of keys in the subtree rooted by c_1 , and let $S(2, x)$ be defined analogously for the right. In addition we impose the constraint all keys in $S(1, x) \leq x$ and all keys in $S(2, x) \geq x$.

Ternary Search Tree: Defined analogously. Given a set $\mathcal{S} := \{k_1, k_2, \dots, k_n\}$ of n integer keys with $k_1 \leq k_2 \leq \dots \leq k_n$: A ternary search tree T_3 is a tree of n ternary search tree nodes. A ternary search tree node can be represented as (x, c_1, c_2, c_3) where $x \in \mathcal{S}$ is a key in \mathcal{S} uniquely represented in T_3 , c_1 is the leftmost child node of x , c_2 is the middle child node of x , and c_3 is the rightmost node of x . Let $S(1, x)$ denote the set of keys in the subtree rooted by c_1 , and let $S(2, x)$ $S(3, x)$ be defined analogously. We impose the constraint that x has to be either greater than or equal to all keys in $S(1, x)$ and $S(2, x)$ and less than or equal to all keys in $S(3, x)$, or x is greater than or equal to keys in $S(1, x)$ and less than or equal to all keys in $S(2, x)$ and $S(3, x)$.

Princeton Computer Science Contest – Fall 2021





Princeton Computer Science Contest – Fall 2021

Thought exercise. Think how accessing a key in a well-balanced ternary search tree takes $O(\log_3 n)$ comparisons.

3 Problem Statements

Given n keys and associated access frequencies F_1, F_2, \dots, F_n , write a program to find out either a optimal binary search tree in $O(n^2)$ -time, or an optimal ternary search tree in $O(n^3)$ -time. *There are two subproblems; solving any part gets you the 15 points credit. Submit both parts using the HackerRank contest platform.*

- Given \vec{F} and n keys $k_1, k_2, \dots, k_n \in \mathbb{Z}$ sorted in increasing order, write a program to optimal cost of an optimal binary search tree, in $O(n^2)$ -time.
 - Input constraint: $1 \leq n \leq 2000$. Use 64-bit signed integer to for your solution so that in case of overflow the behavior will be the same.
 - Running time limit: 1 second.
 - Hint: An $O(n^3)$ solution will not work.
- Given \vec{F} and n keys $k_1, \dots, k_n \in \mathbb{Z}$ sorted in increasing order, write a program that outputs the optimal cost of an optimal ternary search tree, in $O(n^3)$ -time.
 - Input constraint: $1 \leq n \leq 300$. Use 64-bit signed integer to for your solution so that in case of overflow the behavior will be the same.
 - Running time limit: 1 second.
 - Hint: An $O(n^4)$ solution will not work.
- Extra credit opportunity: Email coscon.written.submission@gmail.com and attach a proof of the following math problem. This can potentially get you an additional 3 points max.¹

A cute math problem. We say a function $f : \mathbb{R} \rightarrow \mathbb{R}$ is convex if $\frac{1}{2}(f(x) + f(y)) \leq f(\frac{x+y}{2})$. Let $g : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ be defined as $g(a, b) = f(b - a)$. Show that $g(a, c) + g(b, d) \geq g(b, c) + g(a, d) \iff f$ is convex, for $a \leq b \leq c \leq d$.

Concrete input format descriptions start on the next page.

¹If you have no idea what the inequality about g is, maybe solving (2) will be a little more accessible to you. **This is strictly not necessarily for solving any part of this problem.**

Princeton Computer Science Contest – Fall 2021





Princeton Computer Science Contest – Fall 2021

Part 1 Input/Output Format: Everything is on a single line. The first number read in will be n , the number of nodes. The next n space-separated integers represent F_1, \dots, F_n , respectively. Your program should output a *single* number represent the optimal cost of a BST on these nodes and frequencies.

Sample Input for Part 1:

6 1 4 8 3 6 5

Sample Output for Part 1:

28

Part 2 Input/Output Format: Everything is on a single line. The first number read in will be n , the number of nodes. The next n space-separated integers represent F_1, \dots, F_n , respectively. Your program should output a *single* number represent the optimal cost of a BST on these nodes and frequencies.

First Sample Input for Part 2:

6 1 4 8 3 6 5

First Sample Output for Part 2:

23

Second Sample Input for Part 2:

4 3 8 4 5

Second Sample Output for Part 2:

12

Third Sample Input for Part 2:

4 1 3 4 8

Third Sample Output for Part 2:

9

Princeton Computer Science Contest – Fall 2021

